

open+local
the now of ai



Ben Hur

[open] technology and geospatial consultant



```
ben@sinag-2026:~$ ./intro.sh
```





- > Hi, my name is **Ben**.
- > I'm an **open technology** and **digital autonomy** advocate.
- > I also make a living helping people and organizations **make the most of open data and open technologies** through **BNHR**, **SmartCT**, and the **Civic Literacy Initiative**.





Ben Hur

[open] technology and geospatial consultant








Ben Hur

[open] technology and geospatial consultant

Linux user since 2013

/e/ OS, Lineage OS, and GrapheneOS on mobile



How do you use
AI tools?

How do you use
AI tools?

What
AI tools do you
use?

How do you use
AI tools?

What
AI tools do you
use?

How much do
these **AI tools**
cost?

enshittification



ENSHITTIFICATION

[en-shit-uh-fi-kay-shun]

noun.

The predictable, three-stage lifecycle where platforms evolve from valuable services into useless digital sludge.

(Coined by author and technologist Cory Doctorow, 2022).



Shifting value from users to advertisers, and then to shareholders.






Stage 1: The Bait

The Tactic: Be amazing. Operate at a massive loss. Offer a flawless, ad-free experience.

The Goal: Lock users in.



The Trap: Social lock-in. It's hard to leave when your entire network of **friends** is already there. You value the **connection** more than you hate the platform.



Stage 2: The Switch

The Tactic: Abuse the trapped users to attract business customers.

The Reality: The platform harvests your data, introduces ads, and tweaks algorithms to favor paid content over organic posts from your friends.

The Result: The experience degrades for you, just to make things better for advertisers.




Stage 3: The Squeeze

The Tactic: Abuse everyone. Extract maximum profit for executives and shareholders.

The Reality: Ad prices skyrocket, feeds become a wasteland of sponsored junk.


The Result: Only a homeopathic residue of value remains—just bare minimum functionality to keep you from deleting the app entirely.

Generative AI is shifting from helping us to hustling us.



The exact same blueprint that ruined social networks is being slapped onto Artificial Intelligence.

We are handing over our cognitive labor, creativity, and decision-making to systems built to eventually exploit us.



How AI Holds You Hostage

**Social networks locked in your friends.
AI locks in your data, workflows, and
thoughts.**

you'll own nothing
and be happy

Rented intelligence comes with three hidden taxes.



The Privacy Tax.

Your proprietary codebase, internal documents, and customer data leave your network to train their next model.



The Token Tax.

A variable OpEx nightmare. Every unit test generated and every agent iteration costs real money, punishing scale.



Model Drift.

Unannounced weight updates from the provider mean the prompt that worked perfectly on Monday produces garbage on Tuesday.

The Illusion of Free AI

Dark patterns to hook or confuse you.

Even if you're paying for the AI, your thoughts, private data, and attention will still be the product.

The risk isn't just that AI becomes bad. It's that it becomes good enough to exploit us, but not free enough to empower us.




free+open tech



A different kind of Free.

Free as in freedom.




A free but proprietary (closed source) product is like free food. It's tempting but it won't be able to feed you for life. Plus, it often comes with strings designed to make you dependent.



Free software means freedom—the freedom to have absolute control, ownership, privacy, and agency over your tools.

The Four Essential Freedoms

If software respects your freedom, it must guarantee these four rights:




**The history of open source
is a history of curious young
people building things that
change the world.**

Linus Torvalds was a 21-year-old student when he started
Linux—the open software that now runs the internet.

**Stop being a passive consumer of corporate platforms.
Become an active creator of your own intelligence.**

support local

Breaking the Enshittification Cycle



The five pillars of the local AI value proposition



Privacy

Queries, documents, and conversations never leave the device.




Cost

Zero ongoing API costs; shifting from variable OpEx to fixed CapEx.



Availability

Functions entirely offline with no rate limits or service outages.



Control

Absolute architectural sovereignty. No forced updates or feature removals.



Customization

Unrestricted ability to fine-tune and build custom custom agentic workflows.

Everyday utilities where the edge matches the cloud



Private RAG

Retrieval-Augmented Generation. Chatting with your own local folders of PDFs and notes instantly without uploading them.



Coding Assistance

Replacing cloud-based auto-complete. Using tools like Ollama paired with VS Code (Continue extension) and models like Qwen2.5-Coder.




Rapid Summarization

Handling daily writing, email drafting, and document summarization with lightweight, instantaneous models.

Escaping the token tax eliminates the fear of experimentation

Cloud / Variable Expense



The Token Tax.

Every test run, unit test, and failed agent iteration costs real money.

Example: 1,000 automated workflow tests per day = ~\$300 daily cost.

Local / Fixed Asset



Architectural Sovereignty.

Cost per token drops to the cost of local electricity.

Once the hardware (e.g., Mac Studio or RTX GPU) is purchased, experimentation is functionally free.

8GB
RAM

(The Baseline)

Runs small, highly capable models locally.

Best for: Llama 3.2 3B, Phi-4 Mini.

16GB
RAM

(The Sweet Spot)

Ideal for everyday developer use and reliable reasoning.

Best for: Llama 3.1 8B, Qwen2.5 7B.

32GB+
Unified
Memory


(The Powerhouse)

Unlocks massive reasoning capabilities and complex Retrieval-Augmented Generation (RAG).

Best for: DeepSeek R1 Distills, Mistral-NeMo 12B, and high-parameter multimodal models.

***VRAM (GPU) recommended**

my local ai setup



OLLAMA



Start building with
open models


```
curl -fsSL https://ollama.com/install.sh | sh
```

paste this in terminal, or [download Ollama](#)

The core engine that allows you to run large language models (LLMs) directly on your local machine with ease.

<https://ollama.com>

OPENWEBUI




A feature-rich, user-friendly web interface for interacting with your local AI models.

<https://docs.openwebui.com/>

Proprietary/online alternative: **ChatGPT, Claude, Gemini**

OPENCODE




An open-source coding assistant or model designed to help write, debug, and explain code within your workflow.

<https://opencode.ai/>

Proprietary/online alternative: **Claude Code**

KHOJ



An AI-powered personal search assistant that indexes your notes and files, making your personal knowledge base searchable.

<https://khoj.dev/>

Proprietary/online alternative: **NotebookLM**

build communities



be an advocate

Inform and educate those around you about software freedom, open source, and digital rights.

Be patient and understanding. Not everyone will be as enthusiastic or passionate as you.

Speak up and push back against laws, rules, and regulations that erode our digital freedoms.

Support organizations that seek to preserve and fight for our digital freedoms.

the local open mapping community



From hosting local meetups...

the local open mapping community



Pista ng Mapa 2019



Pista ng Mapa 2022



Pista ng Mapa 2023



State of the Map 2025

...to international conferences

maraming salamat.
humayo kayo at magpaka-open source

